

Lightweight compositional analysis of metagenomes with sourmash gather

This manuscript ([permalink](#)) was automatically generated from [dib-lab/2020-paper-sourmash-gather@41e052e](#) on October 15, 2021.



Authors

- **Luiz Irber**

 [0000-0003-4371-9659](#) ·  [luizirber](#) ·  [luizirber](#)

Graduate Group in Computer Science, UC Davis; Department of Population Health and Reproduction, UC Davis ·
Funded by Grant XXXXXXXX

- **C. Titus Brown**

 [0000-0001-6001-2677](#) ·  [ctb](#)

Department of Population Health and Reproduction, UC Davis

Abstract

The assignment of genomes and taxonomy to metagenome data underlies many microbiome studies. Here we describe two algorithms for compositional analysis of metagenome sequencing data. We first develop a sketching technique, *Scaled MinHash*, that supports Jaccard containment estimation. We implement *Scaled MinHash* in the sourmash software and demonstrate large-scale containment searches of metagenomes using all 700,000 currently available microbial reference genomes. We next frame shotgun metagenome compositional analysis in terms of min-set-cover, i.e. as the problem of finding a minimum collection of reference genomes that “cover” the known portion of a metagenome. We implement a greedy approximate solution that uses *Scaled MinHash* sketches. Finally, we show that by linking genomes to their taxonomic lineages, we can provide a lightweight and precise method for taxonomic classification of metagenome content. sourmash is available as open source under the BSD 3-Clause license at github.com/dib-lab/sourmash/.

Introduction

Shotgun metagenomics samples the DNA sequence content of microbial communities.

Compositional analysis of shotgun metagenome samples has the goal of identifying what reference genomes to use for functional and taxonomic interpretation of metagenome content.

The substantial increase in the number of available reference genomes presents a significant practical obstacle to comprehensive compositional analyses.

Here, we describe a lightweight approach to compositional analysis of shotgun metagenome data. Our approach tackles the selection of appropriate reference genomes and provides a computationally efficient method for taxonomic classification of metagenome data.

We first define *Scaled MinHash*, an extension of MinHash sketching that supports lightweight containment estimation for metagenome datasets using k-mers. We implement *Scaled MinHash* in Python and Rust, and show that it is competitive in accuracy with other containment estimation approaches.

We next frame reference-based metagenome content analysis as a min-set-cov problem, in which we seek the *minimum* number of genomes from the reference database needed to cover the identifiable genomic content of a metagenome. We implement a best-polynomial-time greedy approximation to the min-set-cov problem using *Scaled MinHash*, and show that it recovers a minimum set of reference genomes for the mappable reads in a metagenome.

Finally, we implement a simple taxonomic classification approach on top of min-set-cov, in which we transfer the taxonomy of the genomes from the set cover to the metagenome. We show that this permits precise and lightweight classification of metagenome content across all taxonomic levels.

Results

Scaled MinHash sketches support accurate containment operations

We define the Scaled MinHash on an input domain of k -mers, W , as follows:

$$\mathbf{SCALED}_s(W) = \{ w \leq \frac{H}{s} \mid \forall w \in W \}$$

where H is the largest possible value in the domain of $h(x)$ and $\frac{H}{s}$ is the value in the Scaled MinHash.

The Scaled MinHash is a mix of MinHash and ModHash [1]. It keeps the selection of the smallest elements from MinHash, while using the dynamic size from ModHash to allow containment estimation. However, instead of taking $0 \bmod m$ elements like $\mathbf{MOD}_m(W)$, a Scaled MinHash uses the parameter s to select a subset of W .

Scaled MinHash supports containment estimation with high accuracy and low bias. (Analytic work from David HERE.)

- approximation formula (eqn 13 from overleaf)
- for queries into large sets (large $|A|$), bias factor is low.
- refer to appendix for derivation.

Given a uniform hash function h and $s = m$, the cardinalities of $\mathbf{SCALED}_s(W)$ and $\mathbf{MOD}_m(W)$ converge for large $|W|$. The main difference is the range of possible values in the hash space, since the Scaled MinHash range is contiguous and the ModHash range is not. This permits a variety of convenient operations on the sketches, including iterative downsampling of Scaled MinHash sketches as well as conversion to MinHash sketches.

A Scaled MinHash implementation accurately estimates containment between sets of different sizes

We compare the *Scaled MinHash* method to CMash (*Containment MinHash*) [2] and Mash Screen (*Containment Score*) [3] for containment queries in the dataset from Shakya et al., 2014, a synthetic mock metagenomic bacterial and archaeal community where the reference genomes are largely known [4]. This data set has been used in several methods evaluations [3].




Figure 1: Letter-value plot [hofmann letter-value 2017?] of the differences from containment estimate to ground truth (exact). Each method is evaluated for $k = \{21, 31, 51\}$, except for Mash with $k = 51$, which is unsupported. **A:** Using all 68 reference genomes found in previous articles. **B:** Excluding low coverage genomes identified in previous articles.

All methods are within 1% of the exact containment on average (Figure 1 A), with CMash consistently underestimating the containment for large k and overestimating for small k . Mash Screen with $n = 10000$ has the smallest difference to ground truth for $k = \{21, 31\}$, followed by smol with scaled=1000 and Mash Screen with $n = 1000$.

Figure 1 B shows results with low-coverage and contaminant genomes (as described in [awad evaluating 2017?] and [ondov mash 2019?]) removed from the database. The number of outliers is greatly reduced, with most methods within 1% absolute difference to the ground truth. CMash still has some outliers with up to 8% difference to the ground truth.

CTB todo: * just use (B) benchmark * use sourmash, not smol

CTB questions:

- should we *just* use (B) benchmark?
- should we add sketch sizes in here more explicitly? e.g. number of hashes kept?
- compares well with others
- How much is missed figure; Poisson calculations? => appendix?

A minimum metagenome cover can be efficiently selected from a large database

We next ask: what is the smallest collection of genomes in a database that contains all of the known k-mers in a metagenome? Formally, for a given metagenome M and a reference database D , what is the minimum collection of genomes in D which contain all of the k-mers in the intersection of D and M ? That is, we wish to find the smallest set $\{G_n\}$ of genomes in D such that

$$k(M) \cap k(D) = \bigcup_n \{k(M) \cap k(G_n)\}$$

This is the *minimum set covering* problem, for which there is a polynomial-time approximation (cite). (Provide algorithm here.)

For very large reference databases such as GenBank (which contains over 700,000 microbial genomes as of January 2021) and GTDB (XXX genomes in release XYZ), this is computationally challenging to do exactly. (Estimate total number of k-mers in genbank!) We therefore implemented the algorithm using *Scaled MinHash* sketches to estimate containment, and used an overlap threshold of 100,000 k-mers in order to eliminate genomes with only small overlaps (see Methods).

Table 1: metagenomes and min-set-cov.

data set	genomes >= 100k overlap	min-set-cov
zyzo mock (SRR12324253)	405,839	19
podar mock (SRR606249)	5800	74
p8808mo11 (iHMP)	96,423	99
hu-s1 oil well (SRR1976948)	1235	135

In Table 1, we show the results of running min-set-cov for four metagenomes against genbank - two mock communities (cite cite), one human microbiome data set from iHMP (cite), and one oil well (cite). Our implementation provides estimates for both the *total* number of genomes with substantial overlap to a query genome, and a *minimum list* of genomes that account for k-mers with overlap in the query metagenome (see Methods - prefetch and gather).

We can see that there are many genomes with large overlaps, due to the redundancy of the database. For example, the zyzo mock contains a *Salmonella* genome, and there are over 200,000 *Salmonella* genomes that match to it in Genbank. Likewise, the iHMP dataset contains many XYZ. Since neither the podar mock nor the oil well community contain genomes from species with substantial representation in genbank, they have many fewer total overlaps.

However, regardless of the number of genome with overlap, the estimated *minimum* collection of genomes is always much smaller. In the cases where the k-mers in the metagenome are mostly identified, this is because of database redundancy: e.g. in the case of the zyzo mock, the min-set-cov algorithm chooses only one *Salmonella* genome from the 200,000+ available. However, in the case of

the oil well sample, much of the sample is not represented in the database, suggesting that the small size of the covering set is because many of the species in the metagenome are not in the database.

For the two mock communities where the species compositions are known, the minimum metagenome cover contains genomes belonging to the correct species. (Ask Taylor: what's the best approach to use here, in advance of taxonomic evaluation? Do we want to move CAMI results just below this section?)

Of the 700,000+ genomes in GenBank, 73 were selected by the min-set cov algorithm, and they covered 205.2k (54.8%) of the 374.6k hashes in the metagenome sketch. Importantly, all 64 known genomes were detected, as well as the ones found by mash screen, etc. etc. The unidentified k-mers were all shite. Species analysis confirms etc. etc.

CTB: mention correspondence with depth of sequencing, as well.

CTB: add k-mer overlap statistics to table?

The greedy min-set-cov algorithm progressively classifies the known k-mers

The approximate algorithm used above iteratively subtracts k-mers belonging to the genome with the highest containment count from the metagenome. This results in a progressive classification of the known k-mers in the metagenome to specific genomes. The resulting rank ordered list of genomes can be informative.

Where reference genomes are mostly contained within the metagenome, typical of high-coverage mock communities or communities from which binned genome collections were built, high-rank matches will be ordered by the size of the genome (Figure XXX, top). Lower rank matches reflect situations where portions of reference genomes are either not represented in a metagenome, or there is substantial strain overlap between reference genomes; in the latter case, the core genome content will be "claimed" by the earlier match, and content specific to the strain genome will be represented by late matches (Figure XXX, bottom; compare ...).

Interestingly, the rank-ordered genome classification plot for the iHMP "real" metagenome with a high overall classification rate contains a long tail of strain remnants, while the two mock communities and the real oil-well community do not (Figure XYZ). For the mock communities this is understandable, because these communities were constructed from a low diversity set of known genomes. The difference between the iHMP and environmental metagenome classification curves may be because the oil-well community matches mainly consist of genomes recovered from this specific metagenome, while the iHMP community is matching to a much larger set of possible references (see Table [1](#), column X).

(CTB: matched plot of % genome covered by mapped reads, and ranking)




Figure 2: K-mer decomposition of a metagenome into constituent genomes. A rank ordering by best match first for 30 genomes from the minimum set cover of the synthetic metagenome from Shakya et al., calculated with 700,000 GenBank genomes. The Y axis is labeled with the name of the genome (per NCBI), and the red circles indicates the number of remaining k-mers (estimated with *Scaled MinHash*) shared between each genome and the metagenome. The green x symbols indicate the total number of k-mers shared between each genome and the metagenome, including those already assigned at previous ranks.

Figure 2 shows the results of this algorithm applied to the synthetic metagenome from Shakya et al.

Minimum metagenome covers can accurately estimate taxonomic composition

- CAMI results
- suggests gather/greedy decomposition is pretty good

We evaluated the accuracy of min-set-cov for metagenome decomposition by using benchmarks from the Critical Assessment of Metagenome Interpretation (CAMI) [sczyrba critical 2017?], a community-driven initiative for reproducibly benchmarking metagenomic methods. We used the mouse gut metagenome dataset [meyer tutorial 2020?], in which a simulated mouse gut metagenome (*MGM*) was derived from 791 bacterial and archaeal genomes, representing 8 phyla, 18 classes, 26 orders, 50 families, 157 genera, and 549 species. 64 samples were generated with *CAMISIM*, with 91.8 genomes present on each sample on average. Each sample is 5 GB in size, and both short-read (Illumina) and long-read (PacBio) simulated sequencing data is available.

Since min-set-cov yields only a collection of genomes rather than a species list, we developed an approach for extracting a taxonomic profile from the minimum set cover. We implemented the following procedure for taxonomic profiling: first, we take the known species designation for each genome from the NCBI taxonomy. Then, we use the rank-ordering from the greedy min-set-cov decomposition algorithm to calculate the fraction of the genome remaining in the metagenome after k-mers belonging to higher-rank genomes have been removed. We use this fraction to weight the contribution of the species designation towards the metagenome taxonomy.

(CTB: maybe we need the rank-order decomposition based on k-mers to be in the section before this.)

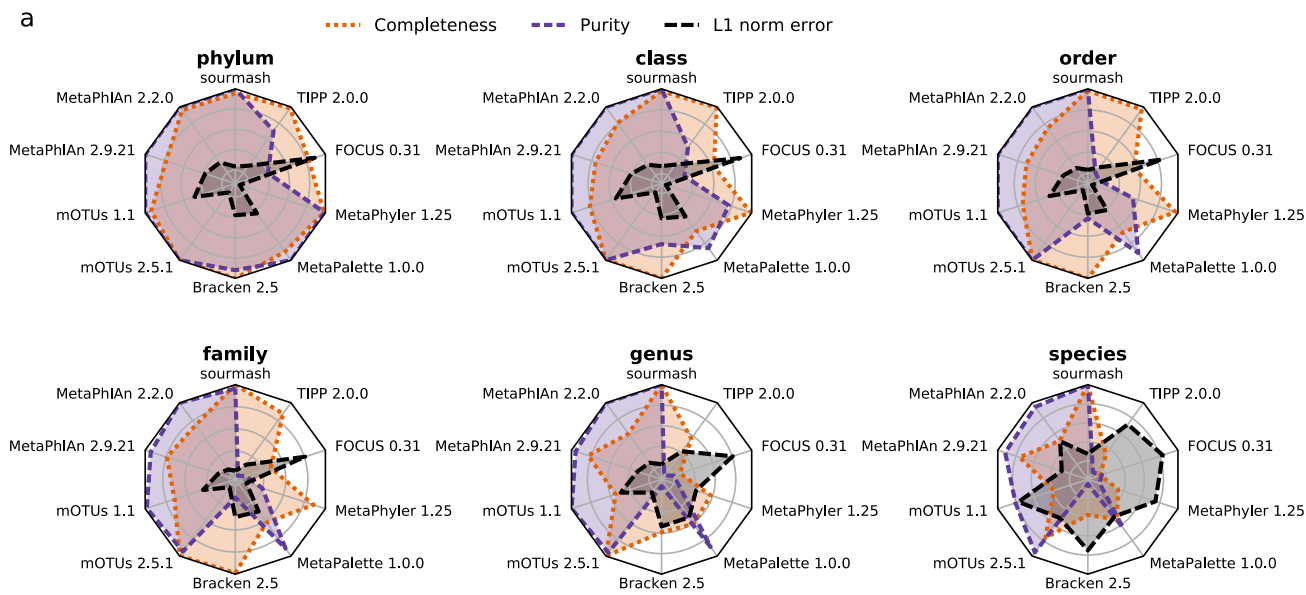


Figure 3: Comparison per taxonomic rank of methods in terms of completeness, purity (1% filtered), and L1 norm.

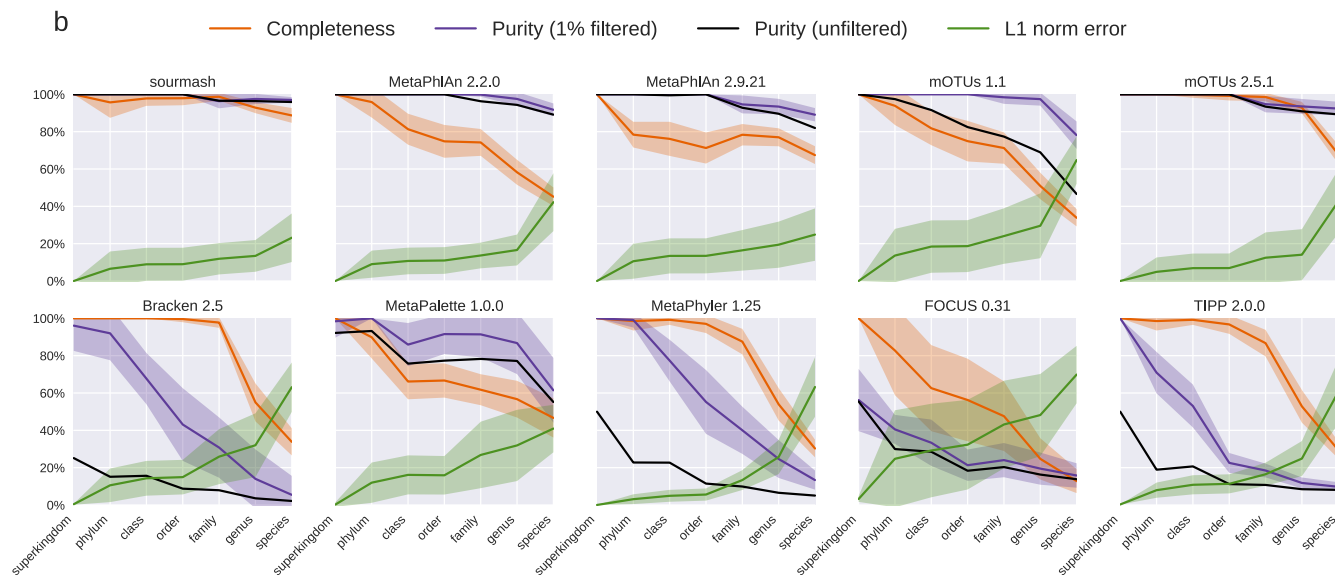


Figure 4: Performance per method at all major taxonomic ranks, with the shaded bands showing the standard deviation of a metric. In **a** and **b**, completeness, purity, and L1 norm error range between 0 and 1. The L1 norm error is normalized to this range and is also known as Bray-Curtis distance. The higher the completeness and purity, and the lower the L1 norm, the better the profiling performance.

C	Completeness	Purity (1% filtered)	L1 norm error	Sum of scores
1st	sourmash (247)	sourmash (179)	mOTUs 2.5.1 (789)	sourmash (1262)
2nd	mOTUs 2.5.1 (416)	MetaPhlAn 2.2.0 (241)	sourmash (836)	mOTUs 2.5.1 (1887)
3rd	Bracken 2.5 (1008)	mOTUs 1.1 (631)	MetaPhlAn 2.9.21 (1401)	MetaPhlAn 2.2.0 (3527)
4th	MetaPhyler 1.25 (1298)	mOTUs 2.5.1 (682)	MetaPhlAn 2.2.0 (1497)	MetaPhlAn 2.9.21 (4349)
5th	TIPP 2.0.0 (1424)	MetaPhlAn 2.9.21 (789)	MetaPhyler 1.25 (1586)	MetaPhyler 1.25 (5148)
6th	MetaPhlAn 2.2.0 (1789)	MetaPalette 1.0.0 (1182)	mOTUs 1.1 (2317)	mOTUs 1.1 (5253)
7th	MetaPhlAn 2.9.21 (2159)	MetaPhyler 1.25 (2264)	TIPP 2.0.0 (2361)	MetaPalette 1.0.0 (5989)
8th	mOTUs 1.1 (2305)	Bracken 2.5 (2881)	MetaPalette 1.0.0 (2390)	Bracken 2.5 (6574)
9th	MetaPalette 1.0.0 (2417)	TIPP 2.0.0 (3361)	Bracken 2.5 (2685)	TIPP 2.0.0 (7146)
10th	FOCUS 0.31 (3424)	FOCUS 0.31 (3764)	FOCUS 0.31 (3894)	FOCUS 0.31 (11082)

Figure 5: Methods rankings and scores obtained for the different metrics over all samples and taxonomic ranks. For score calculation, all metrics were weighted equally.

Figure 3, 4, 5 is an updated version of Figure 6 from [meyer tutorial 2020?] including `sourmash`, comparing 10 different methods for taxonomic profiling and their characteristics at each taxonomic rank. While previous methods show reduced completeness, the ratio of taxa correctly identified in the ground truth, below the genus level, `sourmash` can reach 88.7% completeness at the species level with the highest purity (the ratio of correctly predicted taxa over all predicted taxa) across all methods: 95.9% when filtering predictions below 1% abundance, and 97% for unfiltered results. `sourmash` also has the lowest L1-norm error (the sum of the absolute difference between the true and predicted abundances at a specific taxonomic rank), the highest number of true positives and the lowest number of false positives.

Table 2: Updated Supplementary Table 12 from [meyer tutorial 2020?]. Elapsed (wall clock) time (h:mm) and maximum resident set size (kbytes) of taxonomic profiling methods on the 64 short read samples of the CAMI II mouse gut data set. The best results are shown in bold. Bracken requires to run Kraken, hence the times required to run Bracken and both tools are shown. The taxonomic profilers were run on a computer with an Intel Xeon E5-4650 v4 CPU (virtualized to 16 CPU cores, 1 thread per core) and 512 GB (536.870.912 kbytes) of main memory.

Taxonomic binner	Time (hh:mm)	Memory (kbytes)
MetaPhlAn 2.9.21	18:44	5,139,172
MetaPhlAn 2.2.0	12:30	1,741,304
Bracken 2.5 (only Bracken)	0:01	24,472
Bracken 2.5 (Kraken and Bracken)	3:03	39,439,796
FOCUS 0.31	13:27	5,236,199
CAMIARKQuikr 1.0.0	16:19	27,391,555
mOTUs 1.1	19:50	1,251,296
mOTUs 2.5.1	14:29	3,922,448
MetaPalette 1.0.0	76:49	27,297,132
TIPP 2.0.0	151:01	70,789,939
MetaPhyler 1.25	119:30	2,684,720
sourmash 3.4.0	16:41	5,760,922

When considering resource consumption and running times, `sourmash` used 5.62 GB of memory with an *LCA index* built from the RefSeq snapshot (141,677 genomes) with *scaled* = 10000 and *k* = 51. Each sample took 597 seconds to run (on average), totalling 10 hours and 37 minutes for 64 samples. MetaPhlan 2.9.21 was also executed in the same machine, a workstation with an AMD Ryzen 9 3900X 12-Core CPU running at 3.80 GHz, 64 GB DDR4 2133 MHz of RAM and loading data from an NVMe SSD, in order to compare to previously reported times in Table 2. MetaPhlan took 11 hours and 25 minutes to run for all samples, compared to 18 hours and 44 minutes previously reported, and correcting the `sourmash` running time by this factor it would likely take 16 hours and 41 minutes in the machine used in the original comparison. After correction, `sourmash` has similar runtime and memory consumption to the other best performing tools (*mOTUs* and *MetaPhlAn*), both gene marker and alignment based tools.

Additional points are that `sourmash` is a single-threaded program, so it didn't benefit from the 16 available CPU cores, and it is the only tool that could use the full RefSeq snapshot, while the other tools can only scale to a smaller fraction of it (or need custom databases). The CAMI II RefSeq

snapshot for reference genomes also doesn't include viruses; this benefits `sourmash` because viral *Scaled MinHash* sketches are usually not well supported for containment estimation, since viral sequences require small scaled values to have enough hashes to be reliable.

Notes:

- private database, private taxonomies are easily supported without reindexing.

Minimum metagenome covers provide representative genomes for mapping

K-mer decomposition of metagenomes approximates read mappability

The main point of this section: if you don't like k-mers, read mapping works also. Or, what if the end game is mapping?

Mapping metagenome reads to representative genomes is an intermediate goal in many microbiome analysis pipelines, but mapping approaches struggle with large, redundant databases. One use case for a minimum metagenome cover is to select a small set of representative genomes to be used for mapping. We therefore developed a hybrid selection and mapping pipeline that first rank-orders the minimum metagenome cover by k-mer containment and then iteratively maps reads to those candidate genomes.

CTB: add mapping to this figure.

Overlapping portions of genomes are identified like so (track down the overlaps!)

Overlapping portions of genomes are identified. (Statistics of # k-mers, etc?)

TODO:

- Provide summaries of % k-mers identified/matched, etc.
- CTB: do we want to do this with all k-mers, not just scaled minhash? Or not. *ralstonia* or something? (the one taylor suggested.)

K-mers have been widely used to approximate mapping (citations). To evaluate the accuracy of our k-mer based min-set-cov approach, we identified the ranked set of matching genomes for each of three metagenomes, and then mapped the metagenome reads to the matching genomes with minimap. We then identified reads that mapped to multiple genomes and assigned them to the best-ranked matching genome. This can be viewed as a mapping-based version of min-set-cov.

(CTB note: could also calculate this with mapping, but not against ALL genomes, only against those already found with gather. This highlights the utility of k-mers, since the mapping based approach could not be used against all genomes.)

Figure 6 shows that mapping results generally correspond to gather results. However, they match more closely for synthetic communities than for real communities, especially as gather rank increases. This is likely because in synthetic communities the reference genomes are closer to the actual content of the metagenome, while in real metagenomes we are mapping to imperfect references. CTB: do this for all four?

In particular, both the remnant k-mer and the remnant mappings decrease substantially with increased gather rank. This is because at the higher ranks we are not mapping to all elements in the genome; e.g. in figure XXX, we see that there is a substantial difference in the total number of bases mapped vs the leftover reads from iterative removal. Here only reads that did not map to higher ranked genomes are mapping.

Inspection of the genome taxonomy show that in these situations, we are mapping to subsets of genomes that are the same species or genus as earlier ranked genomes. Figure XYZ compares the best-ranked hash count to the aggregate hash count for the species pangenome; for many species, the aggregate hashes identified for each species in total far outweighs the hashes identified for any one genome.

(belongs in discussion) This suggests that metagenome reads are being mapped to different genomic elements from a species pangenome. While we do not have the resolution to determine this, the most parsimonious interpretation is that the “true” reference genome for the species present in the sample is not in the database, and instead is being cobbled together from core and accessory genome elements in the database.

(Maybe this is where we use *R. gnavus* genomes? Yes - take JUST reads that map to *R. gnavus*, do gather, show what happens x all *gnavus* genomes? Could also do withholding, to show that pangenome elements will usually map one way or another.)

(Show plots with leftover mapping vs all mapping.)

(maybe use sgc here? if so, this would be the last section!)

(CTB: revisit CMash/mash screen papers here to see how they evaluated. Also, maybe mention sgc gbio paper and recovery of new genome.)




Figure 6: Hash-based decomposition of a metagenome into constituent genomes compares well to bases covered by read mapping. The reference genomes are rank ordered along the x axis based on the largest number of hashes from the metagenome specific to that genome, i.e. by order in gather output; hence the number of hashes classified for each genome (orange dots) is monotonically decreasing. The y axis shows absolute number of estimated k-

mers classified to this genome (orange) or total number of bases covered in the reference (blue); the numbers have not been rescaled. Decreases in mapping (green peaks) occur for genomes which are not exact matches to the genomes of the organisms used to build the mock community (cite sherine, mash screen).

Discussion

Scaled MinHash provides efficient compositional queries for large data sets.

Scaled MinHash is an implementation of ModHash that uses the bottom hashing concept from MinHash: all elements in the set to be sketched are hashed, and any hash below a certain fixed boundary value are kept. This fixed boundary value is determined by the desired accuracy for the sketch representation. Unlike MinHash, *Scaled MinHash* supports containment analysis between sets of very different sizes, and here we demonstrate that it can be used efficiently and effectively for compositional analysis of shotgun metagenome data sets with k-mers. In particular, *Scaled MinHash* is competitive in accuracy with extant MinHash-based techniques for containment analysis, while also supporting Jaccard similarity. Footnote: We note that others have also applied the ModHash concept to genomic data; see, for example, Durbin's "modimizer" [8].

Intuitively, Scaled MinHash performs a density sampling at a rate of 1 k -mer per s distinct k -mers seen, where s is the size of the hash space divided by the boundary value used in creating the sketch. This is a kind of lossy compression with a compression ratio of s : that is, for typical values of s used here ($s = 1000$), data sets are reduced in size 1000-fold.

No hash is ever removed from a Scaled MinHash during construction; while this means that sketches grow proportionally to the number of distinct k -mers in the sampled data set, they also support many operations without needing to revisit the original data set. This is in contrast to MinHash, which requires auxiliary data structures for many operations - most especially, containment operations (cite CMash and mash screen). Thus Scaled MinHash sketches serve as distributed compressed indices for the original content for a much broader range of operations than MinHash.

Because *Scaled MinHash* sketches collect all hash values below a fixed threshold, they support streaming analysis of sketches: any operations that used a previously selected value can be cached and updated with newly arriving values. ModHash has similar properties, but this is not the case for MinHash, since after n values are selected any displacement caused by new data can invalidate previous calculations.

Scaled MinHash also directly supports the addition and subtraction of hash values from a sketch, allowing post-processing and filtering without revisiting the original data set. This includes unions and intersections. Although possible for MinHash, in practice this requires oversampling (using a larger n) to account for possibly having less than n values after filtering (the approach taken by Finch [bovee finch: 2018?]).

Another useful operation available on *Scaled MinHash* is *downsampling*: the contiguous value range for Scaled MinHash sketches allow deriving **SCALED** _{s'} (W) sketches for any $s' \geq s$ using only **SCALED** _{s} (W). MinHash and ModHash can also support this operation in limited circumstances, when $n' \leq n$ and m' is a multiple of m . Note also that Scaled MinHash and regular MinHash sketches can be converted between each other when compatible hashing schemes are used, and when (insert math here about boundary values etc.)

Abundance filtering is another extension to MinHash sketches, keeping a count of how many times a value appears in the original data. This allows removing low-abundance values, as implemented in Finch [[bovee finch: 2018?](#)], another MinHash sketching software for genomics. Filtering values that only appear once was implemented in Mash by using a Bloom Filter and only adding values after they were seen once, with later versions also implementing an extra counter array to keep track of counts for each value in the MinHash. These operations can be done in Scaled MinHash without auxiliary data structures.

In exchange for these many conveniences, *Scaled MinHash* sketches have limited sensitivity for small data sets (data set size approx. s) and are only bounded in size by H/s , which is typically quite large. This limited sensitivity may affect the sensitivity of gene- and viral genome-sized queries, but at $s = 1000$ we see comparable accuracy and sketch size to MinHash for bacterial genome comparisons.

(CTB: maybe remove below:)

The consistency of operating in the same data structure also allows further methods to be developed using only *Scaled MinHash* sketches and their features, especially if large collections of *Scaled MinHash* sketches are available. Because Scaled MinHash are collections of hashes, existing k-mer indexing approaches can be applied to the sketches to provide fast database search of these indices.

min-set-cov supports accurate compositional analysis of metagenomes.

Many metagenome content analysis approaches use reference genomes to interpret metagenome content. Here, we frame the computational challenge of discovering the appropriate reference genomes for a set of metagenome reads as a min-set-cov problem, in which we seek a *minimum* set of reference genomes necessary to account for all k-mers shared between the metagenome and the reference database. We show that this can be resolved efficiently for real-world data sets using a greedy algorithm; using *Scaled MinHash*, we provide an approach that scales to 700,000 genomes on current hardware.

The development of a parsimonious list of relevant genomes is convenient in the age of large reference databases with many redundant genomes. Building a parsimonious list of covering genomes from large genome databases is an important feature of the min-set-cov approach. While most extant approaches rely on a short, curated list of non-redundant genomes, ... With Scaled MinHash and sourmash awesomeness, it is now straightforward to routinely search millions of genomes and boil them down to mere dozens of genomes of relevance to a particular metagenome, following which more compute-intensive approaches can be used for detailed analysis. Of course large genome databases may suffer from problems of contamination etc etc but at least this approach gives us the option. (CTB refactor)

This min-set-cov approach for assigning genomes to metagenomes using k-mers differs substantially from extant k-mer and mapping-based approaches. LCA-based approaches such as Kraken assign taxonomy to individual k-mers based on taxonomic lineages in a database, and then use the resulting database of annotated k-mers to assign taxonomy to individual reads or the bulk metagenome. Mapping- and homology-based approaches such as Diamond or @@@ use read mapping to genomes or read alignment to gene sequences to assign taxonomy and function. These approaches typically focus on individual k-mers or reads, in contrast to the greedy min-set-cov approach described here, which looks at the entire collection of reads/k-mers and assigns them in aggregate to the best genome match.

For discussion section:

- LCA is tied to taxonomy, not directly to genomes
- LCA saturates as database sizes grow - more k-mers get pulled up
- in contrast here, we identify combinatorial collections of k-mers in a greedy fashion. this basically means that we pull high-rank/multi-genome k-mers into the largest collection of genome-specific k-mers; we need to evaluate the consequences of this (and do so in the taxonomy section, below).
- also note: LCA chooses discriminatory k-mers in advance, and when databases are updated they must also be updated; here we don't need to do that. We also don't need a step of assigning taxonomy to k-mers.
- it is not clear how important this is computationally in terms of efficiency, given the tradeoffs of the min set cov algorithm, but it should be mentioned.

Unlike Kraken-type approaches, min-set-cov analysis is not tied to taxonomic assignment of genomes; this leads to both computational efficiency in making downstream taxonomic assignments (see discussion below) as well as providing robustness in the face of changing taxonomy.

The greedy algorithm used to determine the minimal list of genomes also lends itself to incremental update with new genomes and supports the use of private databases.

Our comparison of hash-based estimation of containment to mapping results in Figure 6 shows that this approach is an accurate proxy for systematic mapping. In particular, hash-based estimation of containment closely matches actual read mapping performance.

This approach is very dependent on the database. In particular, in many cases the exact reference strains present in the metagenome will not be present in the database. This manifests in two ways in Figure ?? . First, there is a systematic mismatch between the hash content and the mapping content (green line), because mapping software is more permissive in the face of small variants than k-mer-based exact matching. Moreover, many of the lower rank genomes in the plot are from the same species but different *strains* as the higher ranked genomes, suggesting that strain-specific portions of the reference are being utilized for matching at lower ranks. In reality, there will usually be a different mixture of strains in the metagenome than in the reference database. Approaches such as spacegraphcats may help resolve this by adapting old references. [cite?].

Mention weighted cover cc David?

Leftover text:

Our implementation of gather does not currently select the set of smallest genomes, but rather the smallest set of genomes. If there are two genomes with equal containment of the k-mers, it is arbitrary as to which one is chosen.

Note that here we are providing one approach / approximation (Scaled MinHash containment) with one shingling approach (k-mers) to tackle metagenome composition for mapping and taxonomy. The min-set-cover approach could be used with exact containment, and/or with other shingling approaches.

xx can we guess at places where gather would break? One is equivalent containment/different genome sizes.

Any data structure supporting both the *containment* $C(A, B) = \frac{|A \cap B|}{|A|}$ and *remove elements* operations can be used as a query with `gather`. For example, a *set* of the *k*-mer composition of the

query supports element removal, and calculating containment can be done with regular set operations. Approximate membership query (AMQ) sketches like the *Counting Quotient Filter* [[pandey_general-purpose 2017?](#)] can also be used, with the benefit of reduced storage and memory usage. Moreover, the collection of datasets can be implemented with any data structure that can do containment comparisons with the query data structure. Here it can be important to have performant containment searches, since `gather` may run `FindBestContainment` many times.

min-set-cov supports accurate taxonomic classification of metagenome content

Once the min-set-cov approach has identified reference genomes, we can build a taxonomic classifier for metagenome content by simply reporting the taxonomies of the constituent genomes. Our initial taxonomic benchmarking show that this approach is competitive for all metrics across all taxonomic levels.

This approach does not result in the taxonomic saturation caused by the increasing size of large reference databases associated with many other k-mer based methods (Kraken, etc.). As long as every genome in the database possesses a distinct combination of k-mers, the min-set-cov approach can disambiguate reference genomes based on this combination. In practice, our use of *Scaled MinHash* k-mer/hash sampling will limit the resolution of our technique for very closely related genomes, because distinct hashes will not be chosen for them.

One convenient feature of this approach to taxonomic analysis is that new or changed taxonomies can be readily incorporated by assigning them directly to genome identifiers; the majority of the compute is involved in finding the reference genomes, which can have assignments in different taxonomic frameworks. For example, sourmash already supports GTDB natively, and will also support the emerging LINS framework. sourmash can also readily incorporate updates to taxonomies, e.g. frequent updates to the NCBI taxonomy, without requiring expensive reanalysis of the primary metagenome data or even redoing the min-set-cov computation.

Finally, as with the underlying min-set-cov algorithm, it is straightforward to support taxonomic analysis using custom databases and/or custom taxonomic assignments; sourmash already supports this natively.

Algorithm is simple, computational performance is great

The algorithms underlying both *Scaled MinHash* and the greedy min-set-cov solution are simple to describe and straightforward to implement. This increases the likelihood of correct implementation, provides opportunities for independent optimization of data structures, and simplifies interoperability between different implementations.

We provide two implementations with this paper: sourmash, a fully supported open source implementation with command-line, Python and Rust APIs; and smol, a much shorter Rust implementation for demonstration purposes.

sourmash supports large scale data analysis

Taxonomic profiling is fundamentally limited by the availability of reference datasets, even if new reference datasets can be derived from clustering possible organisms based on sequence data in

metagenomes [[milanese microbial 2019?](#)]. The sourmash project provides large scale databases for NCBI and GTDB taxonomies, and supports search of all available genomes.

Limitations of gather

(For *Scaled MinHash*, `gather`, and taxonomy. Move where? Conclusions?)

`gather` as implemented in `sourmash` has the same limitations as *Scaled MinHash* sketches, including reduced sensitivity to small genomes/sequences such as viruses. *Scaled MinHash* sketches don't preserve information about individual sequences, and short sequences using large scaled values have increasingly smaller chances of having any of its k -mers (represented as hashes) contained in the sketch. Because it favors the best containment, larger genomes are also more likely to be chosen first due to their sketches have more elements, and further improvements can take the size of the match in consideration too. Note that this is not necessarily the *similarity* $J(A, B)$ (which takes the size of both A and B), but a different calculation that normalizes the containment considering the size of the match.

`gather` is also a greedy algorithm, choosing the best containment match at each step. Situations where multiple matches are equally well contained or many datasets are very similar to each other can complicate this approach, and additional steps must be taken to disambiguate matches. The availability of abundance counts for each element in the *Scaled MinHash* is not well explored, since the process of *removing elements* from the query doesn't account for them (the element is removed even if the count is much higher than the count in the match). Both the multiple match as well as the abundance counts issues can benefit from existing solutions taken by other methods, like the *species score* (for disambiguation) and *Expectation-Maximization* (for abundance analysis) approaches from Centrifuge [[kim centrifuge 2016?](#)].

(From David Koslicki) Gotchas:

- Lack of sensitivity for small queries
- Potentially large sketch sizes

And a couple other that I've tentatively/mathematically observed:

- The variance of the estimate of $C(A,B) = |AB| / |A|$ appears to also depend on $|A|$, which was somewhat surprising
- The "fixed k-size" problem (which might be able to be overcome with the prefix-lookup data structure, if one sacrifices some accuracy)

Conclusion

- scaled min hash is powerful, with well defined limitations.
- gather is awesome and convenient.
- taxonomy is awesome and overcomes limitations of many current approaches.
- sourmash is robust software that provides a practically usable implementation of these ideas.
- future directions...

Scaled MinHash sketches are simple to implement and analyze, with consistent guarantees for the range of values and subsetting properties when applied to datasets. Containment and similarity operations between *Scaled MinHash* sketches avoid the need to access the original data or more limited representations that only allow membership query, and serve as a proxy for large scale comparisons between hundreds or thousands of datasets.

Small genomes require low scaled values in order to properly estimate containment and similarity, and exact k -mer matching is brittle when considering evolutionarily-diverged organisms. While some of these problems can be overcome in future work, *Scaled MinHash* sketches can serve as a prefilter for more accurate and computationally expensive applications, allowing these methods to be used in larger scales by avoiding processing data that is unlikely to return usable results.

Scaled MinHash sketches are effective basic building blocks for creating a software ecosystem that allow practical applications, including taxonomic classification in metagenomes and large scale indexing and searching in public genomic databases.

Methods

Implementation of Scaled MinHash

We provide two implementations of Scaled MinHash, `smol` and `sourmash`. `smol` is a minimal implementation of *Scaled MinHash* developed to demonstrate the method; it does not include many required features for working with real biological data, but its smaller code base makes it a more readable and concise example of the method. `sourmash` [9] implements features and functionality needed for large scale analyses of real data.

Comparison between CMash, mash screen, and Scaled MinHash.

Experiments use $k = \{21, 31, 51\}$ (except for Mash, which only supports $k \leq 32$). For Mash and CMash they were run with $n = \{1000, 10000\}$ to evaluate the containment estimates when using larger sketches with sizes comparable to the Scaled MinHash sketches with *scaled* = 1000. The truth set is calculated using an exact k -mer counter implemented with a *HashSet* data structure in the Rust programming language [matsakis rust 2014?].

For *Mash Screen* the ratio of hashes matched by total hashes is used instead of the *Containment Score*, since the latter uses a k -mer survival process modeled as a Poisson process first introduced in [fan assembly 2015?] and later used in the *Mash distance* [ondov mash: 2016?] and *Containment score* [ondov mash 2019?] formulations.

MHBT

The *MinHash Bloom Tree* (MHBT) is a variation of the *Sequence Bloom Tree* (SBT) that uses Scaled MinHash sketches as leaf nodes instead of Bloom Filters as in the SBT. The search operation in SBTs is defined as a breadth-first search starting at the root of the tree, using a threshold of the original k -mers in the query to decide when to prune the search. MHBTs use a query Scaled MinHash sketch instead, but keep the same search approach. The threshold of a query Q approach introduced in [solomon fast 2016?] is equivalent to the containment

$$C(Q, S) = \frac{|Q \cap S|}{|S|}$$

described in [broder resemblance 1997?], where S is a Scaled MinHash sketch. For internal nodes n (which are Bloom Filters) the containment of the query Scaled MinHash sketch Q is

$$C(Q, n) = \frac{|\{h \in n \mid \forall h \in Q\}|}{|Q|}$$

as defined by [\[koslicki improving 2019?\]](#) for the *Containment MinHash* to *Bloom Filter* comparison.

MHBTs support both containment and similarity queries. For internal nodes the containment $C(Q, n)$ is used as an upper-bound of the similarity $J(Q, n)$:

$$C(Q, n) \geq J(Q, n) \setminus \frac{|\{Q \cap n\}|}{|\{Q\}|} \geq \frac{|\{Q \cap n\}|}{|\{Q \cup n\}|}$$

since $|Q \cup n| \geq |Q|$. When a leaf node is reached then the similarity $J(Q, S)$ is calculated for the Scaled MinHash sketch S and declared a match if it is above the threshold t . Because the upper-bound is being used, this can lead to extra nodes being checked, but it simplifies implementation and provides better correctness guarantees.

Inverted index

The LCA index in `sourmash` is an inverted index that stores a mapping from hashes in a collection of signatures to a list of IDs for signatures containing the hash. Despite the name, the list of signature IDs is not collapsed to the lowest common ancestor (as in `kraken`), and is calculated as needed by downstream methods using taxonomy information stored separately in the LCA index.

The mapping from hashes to signature IDs in the LCA index is an implicit representation of the original signatures used to build the index, and so returning the signatures is implemented by rebuilding the original signatures on-the-fly. Search in an LCA index matches the k -mers in the query to the list of signatures IDs containing them, using a counter data structure to sort results by number of hashes per signature ID. The rebuilt signatures are then returned as matches based on the signature ID, with containment or similarity to the query calculated against the rebuilt signatures.

`mash screen` [\[ondov mash 2019?\]](#) has a similar index, but it is constructed on-the-fly using the distinct hashes in a sketch collection as keys, and values are counters initially set to zero. As the query is processed, matching hashes have their counts incremented, and after all hashes in the query are processed then all the sketches in the collection are checked in the counters to quantify the containment/similarity of each sketch in the query. The LCA index uses the opposite approach, opting to reconstruct the sketches on-the-fly.

References

1. **On the resemblance and containment of documents**
AZ Broder
Institute of Electrical and Electronics Engineers (IEEE) (2002-11-23) <https://doi.org/fqk7hr>
DOI: [10.1109/sequen.1997.666900](https://doi.org/10.1109/sequen.1997.666900)
2. **IMPROVING MIN HASH VIA THE CONTAINMENT INDEX WITH APPLICATIONS TO METAGENOMIC ANALYSIS**
David Koslicki, Hooman Zabeti
Cold Spring Harbor Laboratory (2017-09-04) <https://doi.org/ghvn6z>
DOI: [10.1101/184150](https://doi.org/10.1101/184150)
3. **Mash Screen: high-throughput sequence containment estimation for genome discovery**
Brian D Ondov, Gabriel J Starrett, Anna Sappington, Aleksandra Kostic, Sergey Koren, Christopher B Buck, Adam M Phillippy
Genome Biology (2019-11-05) <https://doi.org/ghtqmb>
DOI: [10.1186/s13059-019-1841-x](https://doi.org/10.1186/s13059-019-1841-x) · PMID: [31690338](https://pubmed.ncbi.nlm.nih.gov/31690338/) · PMCID: [PMC6833257](https://pubmed.ncbi.nlm.nih.gov/PMC6833257/)
4. **Comparative metagenomic and rRNA microbial diversity characterization using archaeal and bacterial synthetic communities**
Migun Shakya, Christopher Quince, James H Campbell, Zamin K Yang, Christopher W Schadt, Mircea Podar
Environmental Microbiology (2013-06) <https://doi.org/f42ccr>
DOI: [10.1111/1462-2920.12086](https://doi.org/10.1111/1462-2920.12086) · PMID: [23387867](https://pubmed.ncbi.nlm.nih.gov/23387867/) · PMCID: [PMC3665634](https://pubmed.ncbi.nlm.nih.gov/PMC3665634/)
5. **Omega: an Overlap-graph de novo Assembler for Metagenomics**
Bahlul Haider, Tae-Hyuk Ahn, Brian Bushnell, Juanjuan Chai, Alex Copeland, Chongle Pan
Bioinformatics (2014-10) <https://doi.org/f6kt42>
DOI: [10.1093/bioinformatics/btu395](https://doi.org/10.1093/bioinformatics/btu395) · PMID: [24947750](https://pubmed.ncbi.nlm.nih.gov/24947750/)
6. **metaSPAdes: a new versatile metagenomic assembler**
Sergey Nurk, Dmitry Meleshko, Anton Korobeynikov, Pavel A Pevzner
Genome Research (2017-05) <https://doi.org/f97jkv>
DOI: [10.1101/gr.213959.116](https://doi.org/10.1101/gr.213959.116) · PMID: [28298430](https://pubmed.ncbi.nlm.nih.gov/28298430/) · PMCID: [PMC5411777](https://pubmed.ncbi.nlm.nih.gov/PMC5411777/)
7. **Evaluating Metagenome Assembly on a Simple Defined Community with Many Strain Variants**
Sherine Awad, Luiz Irber, CTitus Brown
Cold Spring Harbor Laboratory (2017-07-03) <https://doi.org/ghvn6x>
DOI: [10.1101/155358](https://doi.org/10.1101/155358)
8. **GitHub - richarddurbin/modimizer: a toolset for fast DNA read set matching and assembly using a new type of reduced kmer**
GitHub
<https://github.com/richarddurbin/modimizer>
9. **sourmash: a library for MinHash sketching of DNA**
C Titus Brown, Luiz Irber
The Journal of Open Source Software (2016-09-14) <https://doi.org/ghdrk5>
DOI: [10.21105/joss.00027](https://doi.org/10.21105/joss.00027)

Scaled MinHash sketches support efficient indexing for large-scale containment queries

CTB: Additional points to raise:

- in-memory representation of sketches may be too big (!!), goal here is on disk storage/low minimum memory for “extremely large data” situation.
- Also/in addition, want ability to do incremental loading of things.
- Note we are not talking here about situations where the indices themselves are too big to download.
- I think rename LCA to revindex. Or make up a new name.

We provide two index data structures for rapid estimation of containment in large databases. The first, the MinHash Bloom Tree (MHBT), is a specialization of the Sequence Bloom Tree [solomon fast 2016?], and implements a k -mer aggregative method with explicit representation of datasets based on hierarchical indices. The second is LCA, an inverted index into sketches, a color-aggregative method with implicit representation of the sketches.

We evaluated the MHBT and LCA databases by constructing and searching a GenBank snapshot from July 18, 2020, containing 725,331 assembled genomes (5,282 Archaea, 673,414 Bacteria, 6,601 Fungi 933 Protozoa and 39,101 Viral). MHBT indices were built with *scaled* = 1000, and LCA indices used *scaled* = 10000. Table 3 shows the indexing results for the LCA index, and Table 4 for the MHBT index.

Table 3: Results for LCA indexing, with *scaled* = 10000 and $k = 21$.

Domain	Runtime (s)	Memory (MB)	Size (MB)
Viral	57	33	2
Archaea	58	30	5
Protozoa	231	3	17
Fungi	999	3	65
Bacteria	12,717	857	446

Table 4: Results for MHBT indexing, with *scaled* = 1000, $k = 21$ and internal nodes (Bloom Filters) using 10000 slots for storage.

Domain	Runtime (s)	Memory (MB)	Size (MB)
Viral	126	326	77
Archaea	111	217	100
Protozoa	206	753	302
Fungi	1,161	3,364	1,585
Bacteria	32,576	47,445	24,639

Index sizes are more affected by the number of genomes inserted than the individual *Scaled MinHash* sizes. Despite Protozoan and Fungal *Scaled MinHash* sketches being larger individually, the Bacterial indices are an order of magnitude larger for both indices since they contain two orders of magnitude more genomes.

Comparing between LCA and MHBT index sizes must account for their different scaled parameters, but as shown in Chapter 1 a *Scaled MinHash* with *scaled* = 1000 when downsampled to *scaled* = 10000 is expected to be ten times smaller. Even so, MHBT indices are more than ten times larger than their LCA counterparts, since they store extra caching information (the internal nodes) to avoid loading all the data to memory during search. LCA indices also contain extra data (the list of datasets containing a hash), but this is lower than the storage requirements for the MHBT internal nodes.

We next executed similarity searches on each database using appropriate queries for each domain. All queries were selected from the relevant domain and queried against both MHBT (*scaled* = 1000) and LCA (*scaled* = 10000), for *k* = 21.

Table 5: Running time in seconds for similarity search using LCA (*scaled* = 10000) and MHBT (*scaled* = 1000) indices.

	Viral	Archaea	Protozoa	Fungi	Bacteria
LCA	1.06	1.42	5.40	26.92	231.26
SBT	1.32	3.77	43.51	244.77	3,185.88

Table 6: Memory consumption in megabytes for similarity search using LCA (*scaled* = 10000) and MHBT (*scaled* = 1000) indices.

	Viral	Archaea	Protozoa	Fungi	Bacteria
LCA	223	240	798	3,274	20,926
SBT	163	125	332	1,656	2,290

Table 5 shows running time for both indices. For small indices (Viral and Archaea) the LCA running time is dominated by loading the index in memory, but for larger indices the cost is amortized due to the faster running times. This situation is clearer for the Bacteria indices, where the LCA search completes in 3 minutes and 51 seconds, while the SBT search takes 54 minutes.

When comparing memory consumption, the situation is reversed. Table 6 shows how the LCA index consistently uses twice the memory for all domains, but for larger indices like Bacteria it uses as much as 10 times the memory as the MHBT index for the same data.

For both runtime and memory consumption, it is worth pointing that the LCA index is a tenth of the data indexed by the MHBT. This highlights the trade-off between speed and memory consumption for both approaches, especially for larger indices.

Notes: * new genomes can be added quickly to SBT.