

Lightweight compositional analysis of metagenomes with sourmash gather

This manuscript ([permalink](#)) was automatically generated from [dib-lab/2020-paper-sourmash-gather@8011aba](#) on October 27, 2020.



Authors

- **Luiz Irber**

 [0000-0003-4371-9659](#) ·  [luizirber](#) ·  [luizirber](#)

Graduate Group in Computer Science, UC Davis; Department of Population Health and Reproduction, UC Davis ·
Funded by Grant XXXXXXXX

- **C. Titus Brown**

 [0000-0001-6001-2677](#) ·  [ctb](#)

Department of Population Health and Reproduction, UC Davis

Abstract

The accurate assignment of genomes and taxonomy to microbiome samples is a significant challenge in microbiome studies. Here we describe *Scaled MinHash*, an extension of MinHash sketching that permits rapid and accurate compositional analysis of shotgun metagenome data sets using k-mers. We implement this approach in the sourmash software, and provide database indexing implementations that support large-scale Jaccard similarity and containment searches across all available reference genomes. We next frame the problem of shotgun metagenome compositional analysis as a min-set-cover challenge and implement a greedy approximate solution in sourmash using *Scaled MinHash*. We show that this approach finds the minimal collection of reference genomes necessary to map all possible metagenome reads, and provides a lightweight and precise method for taxonomic classification of metagenome content. sourmash is freely available under the BSD 3-Clause license at github.com/dib-lab/sourmash/.

Introduction

Shotgun metagenomics measures the sequence content of microbial communities.

Compositional analysis of shotgun metagenome samples addresses the question of what reference genomes should be used for functional and taxonomic interpretation of metagenome content.

Below, we describe a lightweight approach to compositional analysis of shotgun metagenome samples using a k-mer based approach.

We first define *Scaled MinHash*, an extension of MinHash sketching that supports lightweight containment estimation for metagenome datasets using a hash-based k-mer downsampling approach; the tradeoff is that sketch sizes are unbounded. We implement *Scaled MinHash* in Python and Rust, and show that it is competitive in accuracy with other containment estimation approaches. We further provide efficient indexed database implementations that support similarity and containment search across hundreds of thousands of reference genomes.

We next frame reference-based metagenome content analysis as a min-set-cov problem, in which we seek the *minimum* number of genomes in the reference database necessary to cover the known content of a metagenome. We implement a known best-polynomial-time greedy approximate to the min-set-cov problem using *Scaled MinHash*, and show that it recovers a minimum set of genomes to which all of the identifiable reads in the metagenome map.

Finally, we implement a simple taxonomic classification approach on top of min-set-cov, in which we transfer the taxonomy of the genomes from the set cover to the metagenome. We show that this permits precise and lightweight classification of metagenome content across all taxonomic levels.

CTB: do we want to mention the stable hash system of Scaled MinHash for unknowns in this paper, or save it for MAGsearch paper/other stuff?

Leftover text:

Compositional data analysis is the study of the parts of a whole using relative abundances [1]. This is a general problem with applications across many scientific fields [???], and examples in biology include RNA-Seq [???], metatranscriptomics [???], microbiome and metagenomics [???]. Taxonomic profiling is a particular instance of this general problem with the goal of finding the identity and

relative abundance of microbial community elements at a specific taxonomic rank (species, genus, family), especially in metagenomic samples [???].

Existing taxonomic profilers use different methods to solve this problem, including aligning sequences to a reference database [???], using marker genes derived from known organisms from reference databases [???] or coupled with unknown organisms clustered from metagenomes [???], and exact k -mer matching using fixed k and lowest common ancestor (LCA) for resolving k -mer assignments matching multiple taxons from a reference database [???] or variable k and assigning multiple taxons per sequence, with an option to reduce it further to the LCA [???].

Once each sequence (from raw reads or assembled contigs) has a taxonomic assignment, these methods resolve the final identity and abundance for each member of the community by summarizing the assignments to a specific taxonomic rank, Taxonomic profiling is fundamentally limited by the availability of reference datasets to be used for assignments, and reporting what percentage of the sample is unassigned is important to assess results, especially in undercharacterized environments such as oceans and soil.

Results

Scaled MinHash sketches support containment operations

- scaled minhash supports similarity and containment

We define the Scaled MinHash on an input domain of k -mers, W , as follows:

$$\mathbf{SCALED}_s(W) = \{ w \leq \frac{H}{s} \mid \forall w \in W \}$$

where H is the largest possible value in the domain of $h(x)$ and $\frac{H}{s}$ is the value in the Scaled MinHash.

The Scaled MinHash is a mix of MinHash and ModHash ([???] broder). As in the former it keeps the smallest elements, and from the latter it adopts the dynamic size to allow containment estimation. However, instead of taking $0 \bmod m$ elements like $\mathbf{MOD}_m(W)$, a Scaled MinHash uses a parameter s to select a subset of W .

Given a uniform hash function h and $s = m$, the cardinalities of $\mathbf{SCALED}_s(W)$ and $\mathbf{MOD}_m(W)$ converge for large $|W|$. The main difference is the range of possible values in the hash space, since the Scaled MinHash range is contiguous and the ModHash range is not. Figure [???] shows an example comparing MinHash, ModHash and Scaled MinHash with the same parameter value.

Intuitively, Scaled MinHash is performing a density sampling at a rate of 1 k -mer per s k -mers seen.

Scaled MinHash permits Jaccard containment estimation between two data sets directly from the sketch ...

A drawback of Scaled MinHash when compared to regular MinHash sketches is the size: the MinHash parameter s sets an upper bound on the size of the sketch, independently of the size of the original data. Scaled MinHash sketches grow proportionally to the original data cardinality, and in the worst case can have up to $\frac{H}{s}$ items.

Scaled MinHash sketches offer a fixed range of possible hash values, but with reduced sensitivity for small datasets when using larger s (scaled) values. A biological example are viruses: at $s = 2000$ many viruses are too small to consistently have a hashed value selected by the *Scaled MinHash* approach. Other *MinHash* approaches sidestep the problem by using hashing and streaming the query dataset (*Mash Screen*) or loading the query dataset into an approximate query membership data structure (*CMash*) to allow comparisons with the variable range of possible hash values, but both solutions require the original data or a more limited data representation than *Scaled MinHash*. The consistency of operating in the same data structure also allows further methods to be developed using only *Scaled MinHash* sketches and their features, especially if large collections of *Scaled MinHash* sketches are available.

Others have also applied the ModHash concept to genomic data; see, for example, Durbin's "modimizer" [2].

Scaled MinHash accurately estimates containment between sets of different sizes

- compares well with others
- supports large-scale sketching of genbank
- How much is missed figure; Poisson calculations?

We next compare the *Scaled MinHash* method implemented in `smol` to *CMash* (*Containment MinHash*) and *Mash Screen* (*Containment Score*) for containment queries in the Shakya dataset [3], a synthetic mock metagenomic bacterial and archaeal community where the reference genomes are largely known. This data set has been used in several methods evaluations [???]. (CTB add SPADes etc refs.)




Figure 1: Letter-value plot [??] of the differences from containment estimate to ground truth (exact). Each method is evaluated for $k = \{21, 31, 51\}$, except for Mash with $k = 51$, since Mash doesn't support $k > 32$. **A:** Using all 68 reference genomes found in previous articles. **B:** Excluding low coverage genomes identified in previous articles.

All methods are within 1% of the exact containment on average (Figure 1 A), with CMash consistently underestimating the containment for large k and overestimating for small k . Mash Screen with $n = 10000$ has the smallest difference to ground truth for $k = \{21, 31\}$, followed by smol with scaled=1000 and Mash Screen with $n = 1000$.

Figure 1 B shows results with low-coverage and contaminant genomes (as described in [??] and [??]) removed from the database. The number of outliers is greatly reduced, with most methods within 1% absolute difference to the ground truth. CMash still has some outliers with up to 8% difference to the ground truth.

CTB questions:

- runtimes?
- should we *just* use (B) benchmark?
- should we add sketch sizes in here more explicitly? e.g. number of hashes kept?
- Where does “Scaled MinHash sketch sizes across GenBank domains” go? (Chp 01 from Luiz thesis)

Scaled MinHash sketches support efficient indexing for large-scale containment queries

CTB: Additional points to raise:

- in-memory representation of sketches may be too big (!!), goal here is on disk storage/low minimum memory for “extremely large data” situation.
- Also/in addition, want ability to do incremental loading of things.
- Note we are not talking here about situations where the indices themselves are too big to download.
- I think rename LCA to revindex. Or make up a new name.

We provide two index data structures for rapid estimation of containment in large databases. The first, the MinHash Bloom Tree (MHBT), is a specialization of the Sequence Bloom Tree [???], and implements a k -mer aggregative method with explicit representation of datasets based on hierarchical indices. The second is LCA, an inverted index into sketches, a color-aggregative method with implicit representation of the sketches.

We evaluated the MHBT and LCA databases by constructing and searching a GenBank snapshot from July 18, 2020, containing 725,331 assembled genomes (5,282 Archaea, 673,414 Bacteria, 6,601 Fungi 933 Protozoa and 39,101 Viral). MHBT indices were built with *scaled* = 1000, and LCA indices used *scaled* = 10000. Table 1 shows the indexing results for the LCA index, and Table 2 for the MHBT index.

Table 1: Results for LCA indexing, with *scaled* = 10000 and $k = 21$.

Domain	Runtime (s)	Memory (MB)	Size (MB)
Viral	57	33	2
Archaea	58	30	5
Protozoa	231	3	17
Fungi	999	3	65
Bacteria	12,717	857	446

Table 2: Results for MHBT indexing, with *scaled* = 1000, $k = 21$ and internal nodes (Bloom Filters) using 10000 slots for storage.

Domain	Runtime (s)	Memory (MB)	Size (MB)
Viral	126	326	77
Archaea	111	217	100
Protozoa	206	753	302
Fungi	1,161	3,364	1,585
Bacteria	32,576	47,445	24,639

Index sizes are more affected by the number of genomes inserted than the individual *Scaled MinHash* sizes. Despite Protozoan and Fungal *Scaled MinHash* sketches being larger individually, the Bacterial indices are an order of magnitude larger for both indices since they contain two orders of magnitude more genomes.

Comparing between LCA and MHBT index sizes must account for their different scaled parameters, but as shown in Chapter 1 a *Scaled MinHash* with *scaled* = 1000 when downsampled to *scaled* = 10000 is expected to be ten times smaller. Even so, MHBT indices are more than ten times larger than their LCA counterparts, since they store extra caching information (the internal nodes) to avoid loading all the data to memory during search. LCA indices also contain extra data (the list of datasets containing a hash), but this is lower than the storage requirements for the MHBT internal nodes.

We next executed similarity searches on each database using appropriate queries for each domain. All queries were selected from the relevant domain and queried against both MHBT (*scaled* = 1000) and LCA (*scaled* = 10000), for $k = 21$.

Table 3: Running time in seconds for similarity search using LCA (*scaled* = 10000) and MHBT (*scaled* = 1000) indices.

	Viral	Archaea	Protozoa	Fungi	Bacteria
LCA	1.06	1.42	5.40	26.92	231.26
SBT	1.32	3.77	43.51	244.77	3,185.88

Table 4: Memory consumption in megabytes for similarity search using LCA (*scaled* = 10000) and MHBT (*scaled* = 1000) indices.

	Viral	Archaea	Protozoa	Fungi	Bacteria
LCA	223	240	798	3,274	20,926
SBT	163	125	332	1,656	2,290

Table 3 shows running time for both indices. For small indices (Viral and Archaea) the LCA running time is dominated by loading the index in memory, but for larger indices the cost is amortized due to the faster running times. This situation is clearer for the Bacteria indices, where the LCA search completes in 3 minutes and 51 seconds, while the SBT search takes 54 minutes.

When comparing memory consumption, the situation is reversed. Table 4 shows how the LCA index consistently uses twice the memory for all domains, but for larger indices like Bacteria it uses as much as 10 times the memory as the MHBT index for the same data.

For both runtime and memory consumption, it is worth pointing that the LCA index is a tenth of the data indexed by the MHBT. This highlights the trade-off between speed and memory consumption for both approaches, especially for larger indices.

Notes: * new genomes can be added quickly to SBT.

Metagenome sketches can be accurately decomposed into constituent genomes by a greedy algorithm, 'gather'

- compare conceptually vs LCA approaches; combinatorial. do we want to do a benchmark of some kind wrt LCA saturation?

- CTB: kind of convoluted logic below! We go from “genomes” to “k-mers” to “greedy estimation” to “scaled minhash estimation”

We can use *Scaled MinHash* sketches to estimate the minimal collection of genomes that serve as a combined reference for a metagenome. Computationally, we frame this problem as follows: for a given metagenome M , what is the smallest collection of genomes in the database D to which all mappable reads in M will be mapped? That is, it is the smallest collection of genomes that explain the observed reads, absent reads for which there is no reference genome.

If we treat both M and all the genomes in D as collections of k-mers, this problem can be framed as a *minimal set covering* problem: find the smallest set $\{G_n\}$ of genomes in D such that

$$k(M) \cap k(D) = \bigcup_n \{k(M) \cap k(G_n)\}$$

We implement a greedy solution to this problem as follows: **(algorithm goes here)** This is a known best-polynomial-time solution to min-set-cov.

Algorithm 1 describes the `gather` method using a generic operation `FindBestContainment`. An implementation for `FindBestContainment` for a list of datasets is presented in Algorithm 2.

Any data structure supporting both the *containment* $C(A, B) = \frac{|A \cap B|}{|A|}$ and *remove elements* operations can be used as a query with `gather`. For example, a *set* of the k -mer composition of the query supports element removal, and calculating containment can be done with regular set operations. Approximate membership query (AMQ) sketches like the *Counting Quotient Filter* [??] can also be used, with the benefit of reduced storage and memory usage. Moreover, the collection of datasets can be implemented with any data structure that can do containment comparisons with the query data structure. Here it can be important to have performant containment searches, since `gather` may run `FindBestContainment` many times.

Since *Scaled MinHash* supports both containment estimation and element removal, we implemented `gather` on top of *Scaled MinHash* and evaluated its performance in finding the number of reads that should map to genomes in a database. CTB: circle back around to talk about what the point of `gather` is: The value is in finding (a) the correct genomes and (b) the bits that are unique.

We evaluated `gather`'s performance on the Shakya data as used above, against GenBank, and compared the genome containment estimation with read mapping. In Figure 2, we show that `gather` accurately estimates both the multimapped reads and the set of reads that maps uniquely to his genome. (CTB: revisit CMash/mash screen papers here to see how they evaluated. Also, maybe mention sgc gbio paper and recovery of new genome.)




Figure 2: Hash-based decomposition of a metagenome into constituent genomes compares well to bases covered by read mapping. The reference genomes are rank ordered along the x axis based on the largest number of hashes from the metagenome specific to that genome, i.e. by order in gather output; hence the number of hashes classified for each genome (orange dots) is monotonically decreasing. The y axis shows absolute number of estimated k-mers classified to this genome (orange) or total number of bases covered in the reference (blue); the numbers have not been rescaled. Decreases in mapping (green peaks) occur for genomes which are not exact matches to the genomes of the organisms used to build the mock community (cite sherine, mash screen).

Leftover text:

We define a greedy algorithm, `gather`, that uses a top-down approach to decompose *Scaled MinHash* sketches into constituent sketches. Starting from the k -mer composition of the query, `gather` iteratively finds a match in a collection of datasets with the largest *containment* of the query (most elements in common), and creates a new query by *removing elements* in the match from the original query. The process stops when the new query doesn't have any more matches in the collection, or a user-provided minimum detection threshold is reached.

Note – * multiple databases can be searched all at once. * public and private databases supported

Taxonomic profiling based on 'gather' is accurate

- CAMI results
- suggests gather/greedy decomposition is pretty good

Taxonomic profiling in `sourmash` is built as an extra step on top of the `gather` algorithm. `gather` returns assignments to a dataset in a collection, and based on that assignment the extra step associates a taxonomic ID and lineage with that genome. Lineages are then summarized at each taxonomic rank for all of the taxonomies in that data set.

To evaluate the performance of taxonomic profiling, we used the mouse gut metagenome dataset [??] from the Critical Assessment of Metagenome Intepretation (CAMI) [??], a community-driven initiative for reproducibly benchmarking metagenomic methods. The simulated mouse gut metagenome (*MGM*) was derived from 791 bacterial and archaeal genomes, representing 8 phyla, 18

classes, 26 orders, 50 families, 157 genera, and 549 species. 64 samples were generated with *CAMISIM*, with 91.8 genomes present on each sample on average. Each sample is 5 GB in size, and both short-read (Illumina) and long-read (PacBio) sequencing data is available.

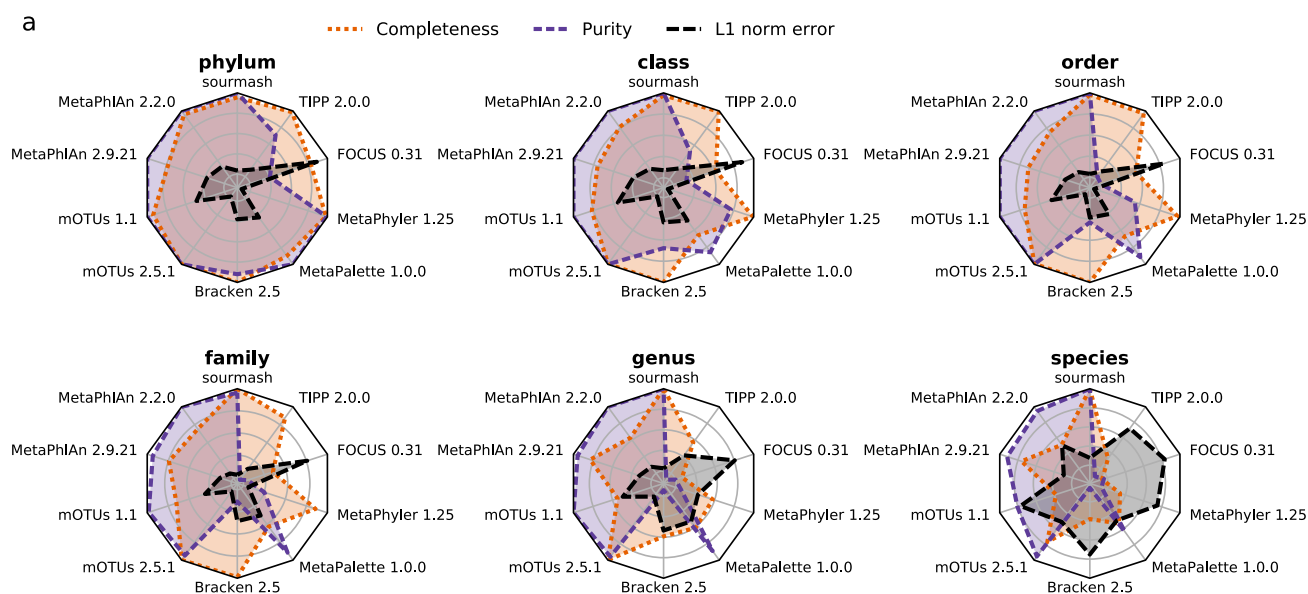


Figure 3: Comparison per taxonomic rank of methods in terms of completeness, purity (1% filtered), and L1 norm.

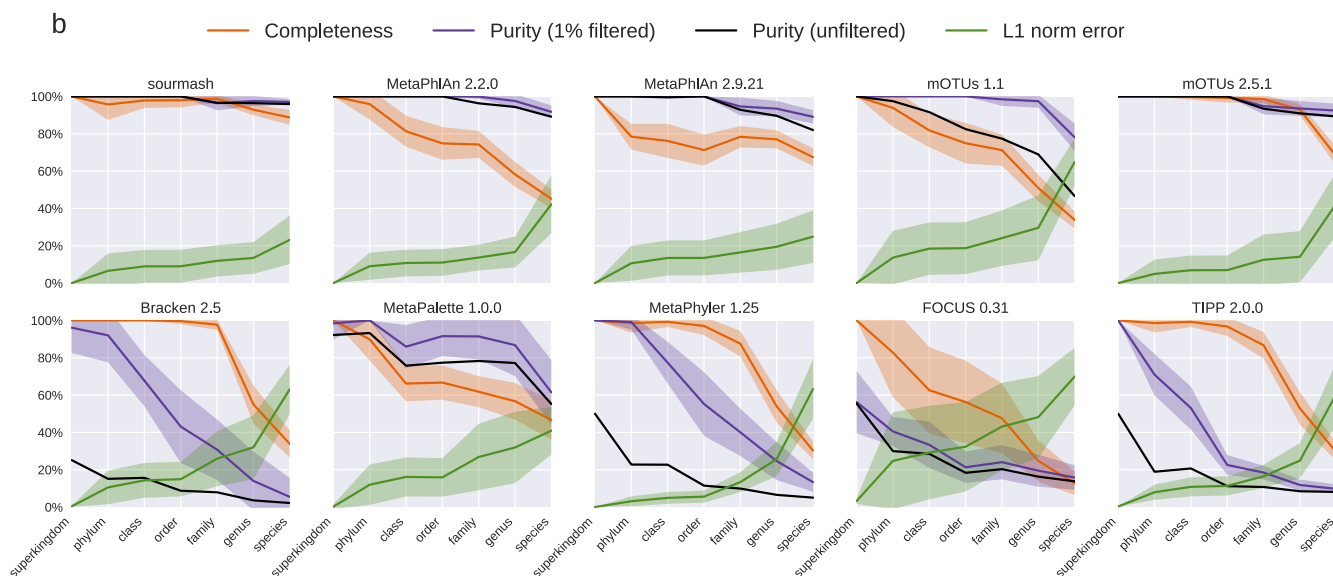


Figure 4: Performance per method at all major taxonomic ranks, with the shaded bands showing the standard deviation of a metric. In **a** and **b**, completeness, purity, and L1 norm error range between 0 and 1. The L1 norm error is normalized to this range and is also known as Bray-Curtis distance. The higher the completeness and purity, and the lower the L1 norm, the better the profiling performance.

C	Completeness	Purity (1% filtered)	L1 norm error	Sum of scores
1st	sourmash (247)	sourmash (179)	mOTUs 2.5.1 (789)	sourmash (1262)
2nd	mOTUs 2.5.1 (416)	MetaPhlAn 2.2.0 (241)	sourmash (836)	mOTUs 2.5.1 (1887)
3rd	Bracken 2.5 (1008)	mOTUs 1.1 (631)	MetaPhlAn 2.9.21 (1401)	MetaPhlAn 2.2.0 (3527)
4th	MetaPhyler 1.25 (1298)	mOTUs 2.5.1 (682)	MetaPhlAn 2.2.0 (1497)	MetaPhlAn 2.9.21 (4349)
5th	TIPP 2.0.0 (1424)	MetaPhlAn 2.9.21 (789)	MetaPhyler 1.25 (1586)	MetaPhyler 1.25 (5148)
6th	MetaPhlAn 2.2.0 (1789)	MetaPalette 1.0.0 (1182)	mOTUs 1.1 (2317)	mOTUs 1.1 (5253)
7th	MetaPhlAn 2.9.21 (2159)	MetaPhyler 1.25 (2264)	TIPP 2.0.0 (2361)	MetaPalette 1.0.0 (5989)
8th	mOTUs 1.1 (2305)	Bracken 2.5 (2881)	MetaPalette 1.0.0 (2390)	Bracken 2.5 (6574)
9th	MetaPalette 1.0.0 (2417)	TIPP 2.0.0 (3361)	Bracken 2.5 (2685)	TIPP 2.0.0 (7146)
10th	FOCUS 0.31 (3424)	FOCUS 0.31 (3764)	FOCUS 0.31 (3894)	FOCUS 0.31 (11082)

Figure 5: Methods rankings and scores obtained for the different metrics over all samples and taxonomic ranks. For score calculation, all metrics were weighted equally.

Figure 3, 4, 5 is an updated version of Figure 6 from [??] including `sourmash`, comparing 10 different methods for taxonomic profiling and their characteristics at each taxonomic rank. While previous methods show reduced completeness, the ratio of taxa correctly identified in the ground truth, below the genus level, `sourmash` can reach 88.7% completeness at the species level with the highest purity (the ratio of correctly predicted taxa over all predicted taxa) across all methods: 95.9% when filtering predictions below 1% abundance, and 97% for unfiltered results. `sourmash` also has the lowest L1-norm error (the sum of the absolute difference between the true and predicted abundances at a specific taxonomic rank), the highest number of true positives and the lowest number of false positives.

Table 5: Updated Supplementary Table 12 from [??]. Elapsed (wall clock) time (h:mm) and maximum resident set size (kbytes) of taxonomic profiling methods on the 64 short read samples of the CAMI II mouse gut data set. The best results are shown in bold. Bracken requires to run Kraken, hence the times required to run Bracken and both tools are shown. The taxonomic profilers were run on a computer with an Intel Xeon E5-4650 v4 CPU (virtualized to 16 CPU cores, 1 thread per core) and 512 GB (536.870.912 kbytes) of main memory.

Taxonomic binner	Time (hh:mm)	Memory (kbytes)
MetaPhlAn 2.9.21	18:44	5,139,172
MetaPhlAn 2.2.0	12:30	1,741,304
Bracken 2.5 (only Bracken)	0:01	24,472
Bracken 2.5 (Kraken and Bracken)	3:03	39,439,796
FOCUS 0.31	13:27	5,236,199
CAMIARKQuikr 1.0.0	16:19	27,391,555
mOTUs 1.1	19:50	1,251,296
mOTUs 2.5.1	14:29	3,922,448
MetaPalette 1.0.0	76:49	27,297,132
TIPP 2.0.0	151:01	70,789,939
MetaPhyler 1.25	119:30	2,684,720
sourmash 3.4.0	16:41	5,760,922

When considering resource consumption and running times, `sourmash` used 5.62 GB of memory with an *LCA index* built from the RefSeq snapshot (141,677 genomes) with *scaled* = 10000 and *k* = 51. Each sample took 597 seconds to run (on average), totalling 10 hours and 37 minutes for 64

samples. MetaPhlan 2.9.21 was also executed in the same machine, a workstation with an AMD Ryzen 9 3900X 12-Core CPU running at 3.80 GHz, 64 GB DDR4 2133 MHz of RAM and loading data from an NVMe SSD, in order to compare to previously reported times in Table 5. MetaPhlan took 11 hours and 25 minutes to run for all samples, compared to 18 hours and 44 minutes previously reported, and correcting the `sourmash` running time by this factor it would likely take 16 hours and 41 minutes in the machine used in the original comparison. After correction, `sourmash` has similar runtime and memory consumption to the other best performing tools (*mOTUs* and *MetaPhlan*), both gene marker and alignment based tools.

Additional points are that `sourmash` is a single-threaded program, so it didn't benefit from the 16 available CPU cores, and it is the only tool that could use the full RefSeq snapshot, while the other tools can only scale to a smaller fraction of it (or need custom databases). The CAMI II RefSeq snapshot for reference genomes also doesn't include viruses; this benefits `sourmash` because viral *Scaled MinHash* sketches are usually not well supported for containment estimation, since viral sequences require small scaled values to have enough hashes to be reliable.

Notes: * private database, private taxonomies are easily supported without reindexing.

Discussion

Scaled MinHash offers benefits, drawbacks vs regular MinHash

TODO: add abundance tracking in to either methods or results; gather bench?

Scaled MinHash is an implementation of ModHash using concepts from MinHashing. *Scaled MinHash* sketches support a variety of options that are convenient for compositional queries - most specifically, containment, but also guarantees on hash occurrence, streaming, hash removal, abundance tracking, and downsampling. Importantly, *Scaled MinHash* sketches can be generated once for large data sets and then used for containment searches after that, unlike CMash and mash screen.

In exchange for these features, *Scaled MinHash* sketches have limited sensitivity for small queries and are bounded in size by H/s , which is usually quite large - so, practically speaking, they grow unbounded with the input size.

Once a Scaled MinHash is calculated there are many operation that can be applied without depending on the original data, saving storage space and allowing scaling analysis to thousands of datasets. Most of these operations are also possible with MinHash and ModHash, with caveats. One example of these operations is : the contiguous value range for Scaled MinHash sketches allow deriving **SCALED** $_{s'}(W)$ sketches for any $s' \geq s$ using only **SCALED** $_s(W)$. MinHash and ModHash can also support this operation, as long as $n' \leq n$ and m' is a multiple of m .

Because Scaled MinHash sketches collect any value below a threshold this also guarantees that once a value is selected it is never discarded. This is useful in streaming contexts: any operations that used a previously selected value can be cached and updated with new arriving values. **MOD** $_m(W)$ has similar properties, but this is not the case for **MIN** $_n(W)$, since after n values are selected any displacement caused by new data can invalidate previous calculations.

Abundance tracking is another extension to MinHash sketches, keeping a count of how many times a value appeared in the original data. This allows filtering for low-abundance values, as implemented in Finch [???], another MinHash sketching software for genomics. Filtering values that only appeared once was implemented before in Mash by using a Bloom Filter and only adding values after they were

seen once, with later versions also implementing an extra counter array to keep track of counts for each value in the MinHash.

TODO: discuss here how abundance tracking in MinHash is not “correct”, because it is not a proper weighted subsample of the data? Note that Scaled MinHash is a proper weighted subsample.

Other operations are adding and subtracting hash values from a Scaled MinHash sketch, allowing post-processing and filtering. Although possible for $\text{MIN}_n(W)$, in practice this requires oversampling (using a larger n) to account for possibly having less than n values after filtering (the approach taken by Finch [???]).

Scaled minhash has limitations vs regular minhash

virus, etc. (could go in first discussion section, but also deserves to be highlighted)

Gather works surprisingly well and matches simple data structures

We next ask, what is a minimal collection of genomes necessary to explain the content of a metagenome? This problem can be framed bioinformatically as asking, for a given metagenome M , what is the smallest collection of genomes in the database D to which all mappable reads in M will be mapped. That is, what is the smallest collection of genomes that explain the observed reads, absent reads for which there is no reference genome?

If we treat both M and all the genomes in D as collections of k-mers, this problem can be framed as a minimal set covering problem: find the smallest set $\{G_n\}$ of genomes in D such that

$$k(M) \cap k(D) = \bigcup_n \{k(M) \cap k(G_n)\}$$

The gather algorithm implements the greedy min-set-cover solution to this problem; this is a known polynomial-time approximation algorithm.

Taxonomic results suggest that this is a pretty good approach.

Our implementation of gather does not currently select the set of smallest genomes, but rather the smallest set of genomes. If there are two genomes with equal containment of the k-mers, it is arbitrary as to which one is chosen.

Note that here we are providing one approach / approximation (Scaled MinHash containment) with one shingling approach (k-mers) to tackle metagenome composition for mapping and taxonomy. The min-set-cover approach could be used with exact containment, and/or with other shingling approaches.

gather is a straightforward algorithm for “decomposing” compositional data. It can take advantage of efficient data structures for containment because it’s “just” k-mers.

gather is a new method for decomposing datasets into its components that outperforms current method when using synthetic datasets with known composition. By leveraging *Scaled MinHash* sketches and efficient indexing data structures it can scale the number of reference datasets used by over an order of magnitude when compared to existing methods.

The `gather` approach differs from previous methods by considering the *co-occurrence* of k -mers between the query and a database sketch as a strong signal that the k -mers originate from that database sketch.

xx can we guess at places where gather would break?

Other containment estimation methods such as *CMash* [???] and *mash screen* [???], can also implement `gather`. Running a search requires access to the original dataset (*mash screen*) for the query, or a Bloom Filter derived from the original dataset (*CMash*), and when the collection of reference sketches is updated the Bloom Filter from *CMash* can be reused, but *mash screen* needs access to the original dataset again.

(Maybe already covered above, or maybe should be moved to “gather can be applied to all the data”?) Since *Scaled MinHash* sketches allow using the sketch directly for `gather`, which are a fraction of the original data in size and also allow enumerating all the elements, an operation not possible with Bloom Filters, they can be stored and reused for large collections of sequencing datasets, including public databases like the Sequence Read Archive [???]. A service that calculate these *Scaled MinHash* sketches and make them available can improve discoverability of these large collections, as well as support future use cases derived from other *Scaled MinHash* features.

Taxonomy results are excellent.

Discuss vs LCA/saturation, slash reference the LCA-has-limits/ k -mers saturate paper

Mix and match taxonomies is easy b/c we anchor to genomes.

Compared to previous taxonomic profiling methods, *Scaled MinHash* can also be seen as a mix of two other approaches: It uses exact k -mer matching and assignment, and the k -mers selected by the MinHashing process are equivalent to implicitly-defined markers. It differs from previous approaches because only a subset of the k -mer composition is used for matching, and traditional gene markers are explicitly chosen due to sequence conservation and low mutation rates, while MinHashing k -mers generates a randomized, but consistent across datasets, set of marker k -mers.

Algorithm is simple, computational performance is great

Performant implementation in sourmasha Python API for data exploration and methods prototyping.

Gather can be applied to all the data.

Taxonomic profiling is fundamentally limited by the availability of reference datasets, even if new reference datasets can be derived from clustering possible organisms based on sequence data in metagenomes [???]. `gather` as implemented in `sourmash` is a method that can scale to increasingly larger collections of datasets due to multiple reasons:

- containment and similarity estimation with *Scaled MinHash* sketches has lower computational requirements than alignment over all reads of a dataset;
- since *Scaled MinHash* sketches use a subset of the k -mer composition, they also scale better than full k -mer composition representations, requiring less space and reducing the number of elements to be computed;

- querying multiple databases can be done independently, avoiding the need to merge, update or reprocess databases when new datasets are available. A new database with the new datasets can be constructed and queried together with previous ones.

Taxonomic profiling in `sourmash` is implemented as an extra step on top of `gather` results. Because these steps are independent of the dataset assignment that `gather` generates, updates to the taxonomy don't require re-executing `gather`, since the taxonomic information can be derived from the same dataset identifier (but potentially with a new associated taxonomic ID). This allows using new taxonomies derived from the same underlying datasets [???], as well as updates to the original taxonomy used before.

Despite improvements to standardization and reproducibility of previous analysis, benchmarking taxonomic profiling tools is still challenging, since tools can generate their reference databases from multiple sources and choosing only one source can bias or make it impossible to evaluate them properly. This is especially true for real metagenomic datasets derived from samples collected from soil and marine environments, where the number of unknown organisms is frequently larger than those contained in reference databases. With the advent of metagenome-assembled genomes (MAGs) there are more resources available for usage as reference datasets, even if they are usually incomplete or draft quality. `sourmash` is well positioned to include these new references to taxonomic profiling given the minimal requirements (a *Scaled MinHash* sketch of the original dataset) and support for indexing hundreds of thousands of datasets.

Limitations of `gather`

`gather` as implemented in `sourmash` has the same limitations as *Scaled MinHash* sketches, including reduced sensitivity to small genomes/sequences such as viruses. *Scaled MinHash* sketches don't preserve information about individual sequences, and short sequences using large scaled values have increasingly smaller chances of having any of its k -mers (represented as hashes) contained in the sketch. Because it favors the best containment, larger genomes are also more likely to be chosen first due to their sketches have more elements, and further improvements can take the size of the match in consideration too. Note that this is not necessarily the *similarity* $J(A, B)$ (which takes the size of both A and B), but a different calculation that normalizes the containment considering the size of the match.

`gather` is also a greedy algorithm, choosing the best containment match at each step. Situations where multiple matches are equally well contained or many datasets are very similar to each other can complicate this approach, and additional steps must be taken to disambiguate matches. The availability of abundance counts for each element in the *Scaled MinHash* is not well explored, since the process of *removing elements* from the query doesn't account for them (the element is removed even if the count is much higher than the count in the match). Both the multiple match as well as the abundance counts issues can benefit from existing solutions taken by other methods, like the *species score* (for disambiguation) and *Expectation-Maximization* (for abundance analysis) approaches from Centrifuge [???].

Future directions for `gather`

In this chapter `gather` is described in terms of taxonomic profiling of metagenomes. That is one application of the algorithm, but it can be applied to other biological problems too. If the query is a genome instead of a metagenome, `gather` can be used to detect possible contamination in the assembled genome by using a collection of genomes and removing the query genome from it (if it is present). This allows finding matches that contain the query genome and evaluating if they agree at specific taxonomic rank, and in case of large divergence (two different phyla are found, for example) it

is likely to indicate that the query genome contains sequences from different organisms. This is especially useful for quality control and validation of metagenome-assembled genomes (MAGs), genomes assembled from reads binned and clustered from metagenomes, as well as a verification during submission of new assembled genomes to public genomic databases like GenBank.

Improvements to the calculation of *Scaled MinHash* sketches can also improve the taxonomic profiling use case. Exact k -mer matching is limited in phylogenetically distant organisms, since small nucleotide differences lead to distinct k -mers, breaking homology assumptions. Different approaches for converting the datasets into a set to be hashed (*shingling*) than computing the nucleotide k -mer composition, such as spaced k -mers [??] and minimizers [??] and alternative encodings for the nucleotides using 6-frame translation to amino acid [??] or other reduced alphabets [??], can allow comparisons on longer evolutionary distances and so improve taxonomic profiling by increasing the sensitivity of the containment estimation. These improvements don't fundamentally change the `gather` method, since it would still be based on the same *containment* and *remove element* operations, but show how `gather` works as a more general method that can leverage characteristics from different building blocks and explore new or improved use cases.

`gather` is a new method for decomposing datasets into its components with application in biological sequencing data analysis (taxonomic profiling) that can scale to hundreds of thousands of reference datasets with computational resources requirements that are accessible to a large number of users when used in conjunction with *Scaled MinHash* sketches and efficient indices such as *LCA* and *MHBT*. It outperforms current methods in community-developed benchmarks, and opens the way for new methods that explore a top-down approach for profiling microbial communities, including further refinements that can resolve larger evolutionary distances and also speed up the method computationally.

XXX SBT and LCA indices

Scaled MinHash sketches are fundamentally a subset of the k -mer composition of a dataset, and so any of the techniques described in [??] are potential candidates for improving current indices or implementing new ones. The MHBT index can be improved by using more efficient representations for the internal nodes [??] and constructing the MHBT by clustering [??], and the LCA index can use more efficient storage of the list of signature IDs by representing the list as colors [??]. The memory consumption of the LCA index can also be tackled by implementing it in external memory using memory-mapped files, letting the operating system cache and unload pages as needed.

Current indices are also single-threaded, and don't benefit from multicore systems. Both indices can be used in parallel by loading as read-only and sharing for multiple searches, but it is also possible to explore parallelization for single queries by partitioning the LCA and assigning each partition to a thread, as well as using a work-stealing thread pool for expanding the search frontier in the MHBT in parallel. In any case, the current implementations serve as a baseline for future scalability and can be used to guide optimization and avoid extraneous overhead and common failings of such projects [??].

Database types work well

"online" approaches

`sourmash gather`, the command-line interface that adds further user experience improvements to the API level, also allows passing multiple indices to be searched, providing incremental support for rerunning with additional data without having to recompute, merge or update the original databases.

Some limitations of gather and database types (equal results can be hard to detect efficiently with current SBT implementation)

The Linear index is appropriate for operations that must check every signature, since it doesn't have any indexing overhead. An example is building a distance matrix for comparing signatures all-against-all. Search operations greatly benefit from extra indexing structure. The MHBT index and k -mer aggregative methods in general are appropriate for searches with query thresholds, like searching for similarity or containment of a query in a collection of datasets. The LCA index and color aggregative methods are appropriate for querying which datasets contain a specific query k -mer.

As implemented in sourmash, the MHBT index is more memory efficient because the data can stay in external memory and only the tree structure for the index need to be loaded in main memory, and data for the datasets and internal nodes can be loaded and unloaded on demand. The LCA index must be loaded in main memory before it can be used, but once it is loaded it is faster, especially for operations that need to summarize k -mer assignments or require repeated searches.

Due to these characteristics, and if memory usage is not a concern, then the LCA index is the most appropriate choice since it is faster. The MHBT index is currently recommended for situations where memory is limited, such as with smaller scaled values ($s \leq 2000$) that increase the size of signatures, or when there are a large number (hundreds of thousands or more) of datasets to index.

Converting between indices

Both MHBT and LCA index can recover the original sketch collection. In the MHBT case, it outputs all the leaf nodes. In the LCA index, it reconstruct each sketch from the hash-to-dataset-ID mapping. This allows trade-offs between storage efficiency, distribution, updating and query performance.

Because both are able to return the original sketch collection, it is also possible to convert one index into the other.

gather Conclusion

Scaled MinHash sketches allow scaling analysis to thousands of datasets, but efficiently searching and sharing them can benefit from data structures that index and optimize these use cases. This chapter introduces an index abstraction that can be trivially implementing using a list of sketches (*Linear index*) and more advanced implementations based on inverted indices (*LCA index*) and hierarchical indices (*MHBT*) providing options for fast and memory-efficient operations, as well as making it easier to share and analyze collections of sketches. All these functionalities are implemented in sourmash.

Limitations and future directions

(For *Scaled MinHash*, `gather`, and taxonomy. Move where? Conclusions?)

(From David Koslicki) Gotchas:

- Lack of sensitivity for small queries
- Potentially large sketch sizes

And a couple other that I've tentatively/mathematically observed:

- The variance of the estimate of $C(A,B) = |AB| / |A|$ appears to also depend on $|A|$, which was somewhat surprising
- The “fixed k -size” problem (which might be able to be overcome with the prefix-lookup data structure, if one sacrifices some accuracy)

Scaled MinHash sketches are fundamentally a subset of the k -mer composition of a dataset, and so any of the techniques described in [??] are potential candidates for improving current indices or implementing new ones. The MHBT index can be improved by using more efficient representations for the internal nodes [??] and constructing the MHBT by clustering [??], and the LCA index can use more efficient storage of the list of signature IDs by representing the list as colors [??]. The memory consumption of the LCA index can also be tackled by implementing it in external memory using memory-mapped files, letting the operating system cache and unload pages as needed.

Current indices are also single-threaded, and don’t benefit from multicore systems. Both indices can be used in parallel by loading as read-only and sharing for multiple searches, but it is also possible to explore parallelization for single queries by partitioning the LCA and assigning each partition to a thread, as well as using a work-stealing thread pool for expanding the search frontier in the MHBT in parallel. In any case, the current implementations serve as a baseline for future scalability and can be used to guide optimization and avoid extraneous overhead and common failings of such projects [??].

Scaled MinHash sketches allow scaling analysis to thousands of datasets, but efficiently searching and sharing them can benefit from data structures that index and optimize these use cases. This chapter introduces an index abstraction that can be trivially implemented using a list of sketches (*Linear index*) and more advanced implementations based on inverted indices (*LCA index*) and hierarchical indices (*MHBT*) providing options for fast and memory-efficient operations, as well as making it easier to share and analyze collections of sketches. All these functionalities are implemented in `sourmash`, a software package exposing these features as a command-line program as well as a Python API for data exploration and methods prototyping.

These indices also serve as another set of building blocks for constructing more advanced methods for solving other relevant biological problems like taxonomic profiling, described in Chapter 3, and approaches for increasing the resilience and shareability of biological sequencing data, described in Chapter 5.

Conclusion

Scaled MinHash sketches are simple to implement and analyze, with consistent guarantees for the range of values and subsetting properties when applied to datasets. Containment and similarity operations between *Scaled MinHash* sketches avoid the need to access the original data or more limited representations that only allow membership query, and serve as a proxy for large scale comparisons between hundreds or thousands of datasets.

Small genomes require low scaled values in order to properly estimate containment and similarity, and exact k -mer matching is brittle when considering evolutionarily-diverged organisms. While some of these problems can be overcome in future work, *Scaled MinHash* sketches can serve as a prefilter for more accurate and computationally expensive applications, allowing these methods to be used in larger scales by avoiding processing data that is unlikely to return usable results.

Scaled MinHash sketches are effective basic building blocks for creating a software ecosystem that allow practical applications, including taxonomic classification in metagenomes and large scale indexing and searching in public genomic databases.

Methods

Implementation of Scaled MinHash

We provide two implementations of Scaled MinHash, `smol` and `sourmash`. `smol` is a minimal implementation of *Scaled MinHash* developed to demonstrate the method; it does not include many required features for working with real biological data, but its smaller code base makes it a more readable and concise example of the method. `sourmash` [4] implements features and functionality needed for large scale analyses of real data.

Comparison between CMash, mash screen, and Scaled MinHash.

Experiments use $k = \{21, 31, 51\}$ (except for Mash, which only supports $k \leq 32$). For Mash and CMash they were run with $n = \{1000, 10000\}$ to evaluate the containment estimates when using larger sketches with sizes comparable to the Scaled MinHash sketches with *scaled* = 1000. The truth set is calculated using an exact k -mer counter implemented with a *HashSet* data structure in the Rust programming language [???].

For *Mash Screen* the ratio of hashes matched by total hashes is used instead of the *Containment Score*, since the latter uses a k -mer survival process modeled as a Poisson process first introduced in [???] and later used in the *Mash distance* [???] and *Containment score* [???] formulations.

MHBT

The *MinHash Bloom Tree* (MHBT) is a variation of the *Sequence Bloom Tree* (SBT) that uses Scaled MinHash sketches as leaf nodes instead of Bloom Filters as in the SBT. The search operation in SBTs is defined as a breadth-first search starting at the root of the tree, using a threshold of the original k -mers in the query to decide when to prune the search. MHBTs use a query Scaled MinHash sketch instead, but keep the same search approach. The threshold of a query Q approach introduced in [???] is equivalent to the containment

$$C(Q, S) = \frac{|Q \cap S|}{|S|}$$

described in [???], where S is a Scaled MinHash sketch. For internal nodes n (which are Bloom Filters) the containment of the query Scaled MinHash sketch Q is

$$C(Q, n) = \frac{|\{h \in n \mid \forall h \in Q\}|}{|Q|}$$

as defined by [???] for the *Containment MinHash* to *Bloom Filter* comparison.

MHBTs support both containment and similarity queries. For internal nodes the containment $C(Q, n)$ is used as an upper-bound of the similarity $J(Q, n)$:

$$C(Q, n) \geq J(Q, n) \iff \frac{|\{Q \cap n\}|}{|Q|} \geq \frac{|\{Q \cap n\}|}{|Q \cup n|}$$

since $|Q \cup n| \geq |Q|$. When a leaf node is reached then the similarity $J(Q, S)$ is calculated for the Scaled MinHash sketch S and declared a match if it is above the threshold t . Because the upper-bound is being used, this can lead to extra nodes being checked, but it simplifies implementation and provides better correctness guarantees.

Inverted index

The LCA index in `sourmash` is an inverted index that stores a mapping from hashes in a collection of signatures to a list of IDs for signatures containing the hash. Despite the name, the list of signature IDs is not collapsed to the lowest common ancestor (as in `kraken`), and is calculated as needed by downstream methods using taxonomy information stored separately in the LCA index.

The mapping from hashes to signature IDs in the LCA index is an implicit representation of the original signatures used to build the index, and so returning the signatures is implemented by rebuilding the original signatures on-the-fly. Search in an LCA index matches the k -mers in the query to the list of signatures IDs containing them, using a counter data structure to sort results by number of hashes per signature ID. The rebuilt signatures are then returned as matches based on the signature ID, with containment or similarity to the query calculated against the rebuilt signatures.

`mash screen [??]` has a similar index, but it is constructed on-the-fly using the distinct hashes in a sketch collection as keys, and values are counters initially set to zero. As the query is processed, matching hashes have their counts incremented, and after all hashes in the query are processed then all the sketches in the collection are checked in the counters to quantify the containment/similarity of each sketch in the query. The LCA index uses the opposite approach, opting to reconstruct the sketches on-the-fly.

References

1. **The Statistical Analysis of Compositional Data**

J. Aitchison

Journal of the Royal Statistical Society: Series B (Methodological) (1982-01) <https://doi.org/gfs7gn>

DOI: [10.1111/j.2517-6161.1982.tb01195.x](https://doi.org/10.1111/j.2517-6161.1982.tb01195.x)

2. **richarddurbin/modimizer**

Richard Durbin

(2020-09-03) <https://github.com/richarddurbin/modimizer>

3. **Comparative metagenomic and rRNA microbial diversity characterization using archaeal and bacterial synthetic communities**

Migun Shakyia, Christopher Quince, James H. Campbell, Zamin K. Yang, Christopher W. Schadt, Mircea Podar

Environmental Microbiology (2013-06) <https://doi.org/f42ccr>

DOI: [10.1111/1462-2920.12086](https://doi.org/10.1111/1462-2920.12086) · PMID: [23387867](https://pubmed.ncbi.nlm.nih.gov/23387867/) · PMCID: [PMC3665634](https://pubmed.ncbi.nlm.nih.gov/PMC3665634/)

4. **sourmash: a library for MinHash sketching of DNA**

C. Titus Brown, Luiz Irber

The Journal of Open Source Software (2016-09-14) <https://doi.org/ghdrk5>

DOI: [10.21105/joss.00027](https://doi.org/10.21105/joss.00027)